

# Optimisation of Parallel and Dense Matrix Multiplication in Scientific Computing

Mara Pareja del Pino<sup>a</sup>

<sup>a</sup>*University of Las Palmas de Gran Canaria Spain*

---

## Abstract

The project explores matrix multiplication techniques with a focus on parallel algorithms for enhanced computational efficiency. The implementation includes a sequential dense matrix multiplier and a parallelized counterpart designed to leverage multicore architectures effectively.

The project's core contribution is the `ParallelMatrixMultiplier`, which intelligently divides matrices into submatrices, distributing computation across multiple threads. This concurrent approach aims to optimize performance, particularly evident in large-scale matrix operations.

Testing strategies, encompassing correctness and efficiency assessments, validate the parallel matrix multiplication against sequential counterparts. The project's flexibility, demonstrated through matrix generation, reading from files, and extensive testing, contributes to its utility in diverse scientific and computational contexts.

This work showcases the potential benefits of parallelization in matrix multiplication, offering insights into scalability and efficiency. It provides a foundation for future research and applications requiring optimized numerical computations.

---

## 1. Introduction

Matrix multiplication is a fundamental operation in various scientific and computational domains, serving as a cornerstone for applications ranging from computer graphics to numerical simulations. As the size of matrices grows, traditional sequential algorithms may struggle to meet the demands of modern computing, necessitating the exploration of parallel computing paradigms. This paper delves into the development and evaluation of parallel matrix multiplication algorithms designed to harness the computational power of multicore architectures.

The project centers around the implementation and analysis of a parallelized matrix multiplication algorithm, aiming to exploit concurrency for accelerated numerical computations. Leveraging the inherent parallelism of matrix operations, particularly evident in the multiplication process, holds promise for significantly enhancing computational efficiency. The study employs a comprehensive testing approach, evaluating the correctness and performance of the parallel algorithm against its sequential counterpart across matrices of varying sizes.

The foundation of the investigation lies in the development of a flexible matrix manipulation framework, accommodating matrix generation, reading from files, and rigorous testing. The core algorithm intelligently decomposes matrices into submatrices, distributing the workload across multiple threads to maximize parallel processing capabilities. This approach is tailored to mitigate the computational challenges posed by larger matrices, offering potential scalability benefits.

Through extensive testing, including performance evaluations and correctness assessments, the paper aims to provide insights into the benefits and trade-offs associated with parallel matrix multiplication. The project's adaptability and utility extend beyond the immediate scope, offering a versatile toolkit for

scientific and computational applications requiring optimized numerical computations. As computational demands continue to grow, the exploration of parallel algorithms for matrix multiplication becomes increasingly imperative, making this research a timely and valuable contribution to the field.

## 2. Problem Statement

The increasing size of matrices in scientific and computational applications poses a challenge to the efficiency of traditional sequential matrix multiplication algorithms. As computational demands escalate, there is a pressing need to explore and develop parallel algorithms that harness the processing power of multicore architectures. The core problem is to strike a balance between algorithmic complexity, load balancing, and effective utilization of parallel resources.

This paper addresses the challenge of optimizing matrix multiplication through the development and evaluation of parallel algorithms. The primary focus is on overcoming the limitations of sequential methodologies and leveraging parallelization to enhance computational efficiency, particularly for large-scale matrix operations. The investigation aims to provide insights into the potential performance gains achievable through parallelization, emphasizing scalability and adaptability across varying matrix sizes.

Additionally, the project endeavors to contribute a flexible framework that extends beyond parallel matrix multiplication, accommodating matrix generation, reading from files, and robust testing methodologies. By addressing the computational challenges associated with matrix multiplication, this research seeks to advance the state-of-the-art in numerical computations and provide a foundation for scalable and efficient parallel algorithms.

### 3. Methodology

#### 3.1. Matrix Representation and Manipulation

We develop a versatile `Matrix` class to represent matrices and facilitate manipulation. Matrix generation allows specification of sizes and fill values for testing, and reading from external files is implemented using the `MatrixReader` class for flexibility with real-world datasets.

#### 3.2. Sequential Matrix Multiplication

A sequential matrix multiplication algorithm (`DenseMatrixMultiplier`) serves as a baseline for performance comparison. Extensive testing ensures correctness against known matrix multiplication results, and the algorithm's performance is evaluated using matrices of varying sizes, recording execution times.

#### 3.3. Parallel Matrix Multiplication

The `ParallelMatrixMultiplier` class introduces a parallelized matrix multiplication algorithm. Matrices are divided into submatrices, distributing computation across multiple threads to exploit parallel processing capabilities. Synchronization mechanisms are implemented to ensure accurate computation and avoid race conditions. Correctness is evaluated through extensive testing against the sequential baseline, and performance is measured and analyzed across varying matrix sizes.

#### 3.4. Performance Comparison

Comprehensive performance comparisons between the sequential and parallel matrix multiplication algorithms are conducted. The `MatrixMultiplicationPerformanceTest` class assesses efficiency for matrices of different sizes, measuring and analyzing execution times to identify scalability trends and performance improvements achieved through parallelization.

#### 3.5. Selected Matrices Testing

The parallel matrix multiplication algorithm is tested against selected matrices of sizes 128, 256, 512, and 1024 read from external files. Correctness is evaluated by comparing results with sequentially computed counterparts. Performance is measured and analyzed for real-world matrices, considering factors such as data size and complexity.

#### 3.6. Flexibility and Utility Evaluation

The flexibility and utility of the developed framework are assessed by testing matrix generation, reading, and manipulation functionalities. The adaptability of project components to diverse scientific and computational applications is explored, and the robustness of the toolkit is validated through extensive testing scenarios, ensuring reliability in various contexts.

#### 3.7. Documentation and Reporting

Implementation details, algorithms, and key design decisions are documented. Experimental results are compiled and analyzed, presented through visualizations and performance metrics. Concise summaries and insights drawn from the methodology are formulated, discussing implications and potential future research directions.

### 4. Project Architecture

The project is organized into several packages, each responsible for specific aspects of the matrix operations and management.

#### 4.1. *org.bigdata.model*

This package contains the `Matrix` class, which serves as the fundamental representation of matrices. The `Matrix` class encapsulates matrix data and provides essential methods for setting and retrieving values, as well as obtaining information about the matrix size.

#### 4.2. *org.bigdata.generator*

The `MatrixGenerator` class, located in this package, is responsible for creating matrices for testing purposes. It includes methods to generate matrices of specified sizes and to save them to files. Additionally, it provides functionality to read matrices from files, offering flexibility in working with both generated and real-world data.

#### 4.3. *org.bigdata.operations*

This package houses the matrix multiplication operations, including both sequential and parallel implementations.

- `DenseMatrixMultiplier`: This class implements the sequential matrix multiplication algorithm. It is responsible for multiplying two matrices in a straightforward, non-parallelized manner.
- `ParallelMatrixMultiplier`: The parallel matrix multiplication is handled by this class. It intelligently divides matrices into submatrices and utilizes multithreading to perform matrix multiplication in parallel. The implementation focuses on optimizing performance for large-scale matrices.
- `SubmatrixesDivider`: Provides utilities for dividing matrices into submatrices. This class is essential for the parallel matrix multiplication algorithm.

#### 4.4. *org.bigdata*

The main package, `org.bigdata`, contains utility classes and the entry point for the application.

- `ThreadCounter`: A simple utility class that determines the number of available threads in the runtime environment.

## 5. Tests Implemented

To validate the correctness and performance of matrix operations, a series of tests were implemented. These tests comprehensively evaluate both sequential and parallel matrix multiplication algorithms.

### 5.1. Sequential Matrix Multiplication Tests

The `DenseMatrixMultiplier` class, responsible for sequential matrix multiplication, underwent extensive testing. Unit tests covered various matrix sizes, comparing expected outcomes against manually calculated results to ensure accuracy.

### 5.2. Parallel Matrix Multiplication Tests

Rigorous testing of the `ParallelMatrixMultiplier` class, which manages parallel matrix multiplication, was conducted. Tests involved matrices of different dimensions, with the parallel implementation compared against the sequential counterpart for accuracy.

### 5.3. Performance Comparison

A systematic comparison of the performance between sequential and parallel matrix multiplication was executed. Matrices of increasing sizes were used, measuring execution times for each approach. The scalability of the parallel implementation was analyzed to identify optimal use cases.

### 5.4. Conclusions of the Tests

The implemented tests provide valuable insights into the performance and efficiency of both sequential and parallel matrix multiplication algorithms. Tests were conducted on matrices of varying sizes, yielding the following results:

#### 5.4.1. Sequential Matrix Multiplication

Execution times for sequential matrix multiplication (`DenseMatrixMultiplier`) are as follows:

- Matrix size 128x128: 17.935 ms
- Matrix size 256x256: 21.7619 ms
- Matrix size 512x512: 196.8893 ms
- Matrix size 1024x1024: 2125.768 ms

The sequential implementation exhibits a clear trend of increasing execution time with larger matrices.

#### 5.4.2. Parallel Matrix Multiplication

Execution times for parallel matrix multiplication (`ParallelMatrixMultiplier`) are as follows:

- Matrix size 128x128: 83.0663 ms
- Matrix size 256x256: 17.6236 ms
- Matrix size 512x512: 44.74 ms
- Matrix size 1024x1024: 329.5146 ms

The parallel implementation demonstrates significant improvements in performance, particularly evident in reduced execution times for larger matrices.




Figure 1: Sequential vs Parallel Matrix Multiplication Times

### 5.4.3. Conclusion of the Tests

The parallel matrix multiplication algorithm demonstrates its effectiveness in performance optimisation, offering a substantial speedup compared to the sequential approach. However, it is worth noting that initially the sequential multiplication is faster than the parallel one, this changes when we reach the matrix size of 256x256 where we can see how the execution times are similar and then the execution time of the sequential multiplication increases well above the parallel one.

The results highlight the scalability of the parallel implementation, with notable time reductions, especially in scenarios involving matrices with large dimensions.

These results confirm the advantages of leveraging parallelisation techniques for matrix operations, especially in scenarios where large-scale computations are predominant.

## 6. Conclusions

The matrix multiplication project yields several noteworthy conclusions:

**Performance Advantages:** The parallel matrix multiplication algorithm (`ParallelMatrixMultiplier`) consistently outperforms the sequential approach (`DenseMatrixMultiplier`) for larger matrices. This performance superiority is particularly evident in scenarios involving extensive computational tasks with substantial matrix dimensions.

**Scalability and Efficiency:** The parallel algorithm's scalability is a notable strength, efficiently handling increased matrix sizes. The distribution of workload across multiple threads highlights the benefits of parallelization, optimizing matrix multiplication operations.

**Utility of Parallelization:** Efficiently leveraging multiple threads proves effective in dividing and conquering the workload, resulting in significant improvements in execution times.

This underscores the relevance of parallelization in scenarios where computational resources can be utilized for performance gains.

**Flexibility in Matrix Operations:** The inclusion of the `MatrixGenerator` class in the `org.bigdata.generator` package adds flexibility by enabling the generation of matrices for testing purposes. Additionally, the capability to read matrices from files enhances the project's versatility, allowing seamless integration with both generated and real-world data.

In summary, the successful implementation of matrix multiplication algorithms, coupled with comprehensive testing and performance evaluations, underscores the project's contribution to the field of parallel computing and numerical operations.

## 7. Future Work

While the matrix multiplication project has achieved success in implementing and optimizing parallel algorithms, several areas offer avenues for future exploration and enhancement:

**Parallelization Strategies:** Investigate and implement alternative parallelization strategies, such as GPU acceleration or distributed computing, to assess their impact on performance and scalability. Exploring diverse parallel architectures may provide additional insights and optimizations.

**Dynamic Workload Balancing:** Enhance the parallel matrix multiplication algorithm with dynamic workload balancing techniques. Implement mechanisms that adaptively distribute submatrices among threads based on runtime characteristics, potentially improving efficiency in varying computational scenarios.

**Algorithmic Improvements:** Explore advanced matrix multiplication algorithms beyond the basic parallel and sequential approaches. Investigate state-of-the-art algorithms, such as Strassen's algorithm, to evaluate their applicability and performance benefits within the project context.

**Integration with Frameworks:** Explore integration possibilities with existing parallel computing frameworks, such as Apache Spark or Apache Flink. Leveraging established frameworks may provide additional optimization opportunities and compatibility with broader distributed computing ecosystems.

**Benchmarking and Optimization:** Conduct extensive benchmarking with diverse matrices and sizes to further refine and optimize the parallel matrix multiplication algorithm. Identify specific matrix characteristics that influence performance and tailor optimizations accordingly.

Continuing research and development in these areas can contribute to the project's evolution, ensuring its relevance and effectiveness in diverse computational contexts.